

Evaluation Run Analysis: eval_claude_100rows

Based on the run data, here's a comprehensive analysis of your evaluation:

Overview

- **Model:** anthropic-claude-sonnet-4-6
- **Samples Evaluated:** 100 rows
- **Status:** ✓ FINISHED (172.1s duration)
- **Error Rate:** 0 (no errors)

Performance Metrics

Latency

- **Mean:** 10.26s per request
- **P90:** 14.99s
- *Insight:* Relatively high latency, likely due to Claude Sonnet 4 processing complexity. Consider Claude Haiku for faster responses if acceptable quality tradeoff.

Cost

- **Average:** \$0.00462 per intake
- *Insight:* ~\$0.46 per 100 intakes. For high-volume use, monitor accumulated costs.

Quality Scores (Scale appears to be 1-5)

Metric	Score	Assessment
Legal Compliance	5.0/5	★ Perfect - strong regulatory adherence
Information Completeness	4.91/5	★ Excellent - comprehensive responses
Professional Tone	4.32/5	✓ Good - appropriate formality
Response Structure	3.65/5	⚠ Needs improvement - organization could be better

Key Finding: While content quality is excellent, response structure is the weakest area. Consider adding formatting guidelines or examples to your prompt.

Classification Accuracy

Task	Score	Status
Practice Area Accuracy	100%	★ Perfect identification
Area Macro F1	1.0	★ Excellent balance
Area Macro Recall	100%	★ No missed practice areas
Category Accuracy	48%	⚠ Major concern
Category Macro Recall	52%	⚠ Missing many categories

Critical Issue: The model perfectly identifies high-level practice areas but struggles with specific category classification. This 48% category accuracy needs attention.

Recommendations

- 1. Address Category Classification** (Priority: High)
 - The 48% category accuracy vs. 100% practice area accuracy suggests the model needs better category definitions or examples
 - Review your category taxonomy - are categories too fine-grained or ambiguous?
 - Add few-shot examples for challenging category distinctions
- 2. Improve Response Structure** (Priority: Medium)
 - Score of 3.65/5 indicates formatting issues
 - Add structured output requirements to your prompt (bullet points, sections, etc.)
 - Consider using JSON output mode for more consistent structure
- 3. Optimize for Production**
 - 10s mean latency may impact user experience
 - Test Claude Haiku 4 as a faster alternative while monitoring quality metrics
 - Consider caching for common queries
- 4. Next Steps**
 - Drill into failed category classifications to identify patterns
 - Run A/B test with improved prompts focusing on category distinctions
 - Set up scheduled scorers to monitor these metrics in production